

Лабораторная работа №5: Ансамбли моделей машинного обучения.

```
In [3]: #Датасет содержит данные о кредитах на покупку электроники, которые были
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score
from sklearn.neural_network import MLPClassifier
from warnings import simplefilter

simplefilter('ignore')
```

```
In [4]: # записываем CSV-файл в объект DataFrame
data = pd.read_csv('credit_train_preprocess.csv', encoding='cp1251', sep=';')
```

```
In [5]: # смотрим на первые пять строк
data.head()
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 170746 entries, 0 to 170745
Data columns (total 39 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   170746 non-null  float64
1   credit_sum                           170746 non-null  float64
2   credit_month                         170746 non-null  int64
3   tariff_id                           170746 non-null  float64
4   score_shk                           170746 non-null  float64
5   monthly_income                      170746 non-null  float64
6   credit_count                        170746 non-null  float64
7   overdue_credit_count               170746 non-null  float64
8   open_account_flg                   170746 non-null  int64
9   gender_F                           170746 non-null  int64
10  gender_M                           170746 non-null  int64
11  job_position_ATP                    170746 non-null  int64
12  job_position_BIS                    170746 non-null  int64
13  job_position_BIU                    170746 non-null  int64
14  job_position_DIR                    170746 non-null  int64
15  job_position_HSK                    170746 non-null  int64
16  job_position_INP                    170746 non-null  int64
17  job_position_INV                    170746 non-null  int64
18  job_position_NOR                    170746 non-null  int64
19  job_position_ONB                    170746 non-null  int64
20  job_position_PNA                    170746 non-null  int64
21  job_position_PNI                    170746 non-null  int64
22  job_position_PNS                    170746 non-null  int64
23  job_position_PNV                    170746 non-null  int64
```

```

24  job_position_SPC      170746 non-null  int64
25  job_position_UMN      170746 non-null  int64
26  job_position_WOI      170746 non-null  int64
27  job_position_WRK      170746 non-null  int64
28  job_position_WRP      170746 non-null  int64
29  education_ACD         170746 non-null  int64
30  education_GRD         170746 non-null  int64
31  education_PGR         170746 non-null  int64
32  education_SCH         170746 non-null  int64
33  education_UGR         170746 non-null  int64
34  marital_status_CIV     170746 non-null  int64
35  marital_status_DIV     170746 non-null  int64
36  marital_status_MAR     170746 non-null  int64
37  marital_status_UNM     170746 non-null  int64
38  marital_status_WID     170746 non-null  int64
dtypes: float64(7), int64(32)
memory usage: 50.8 MB

```

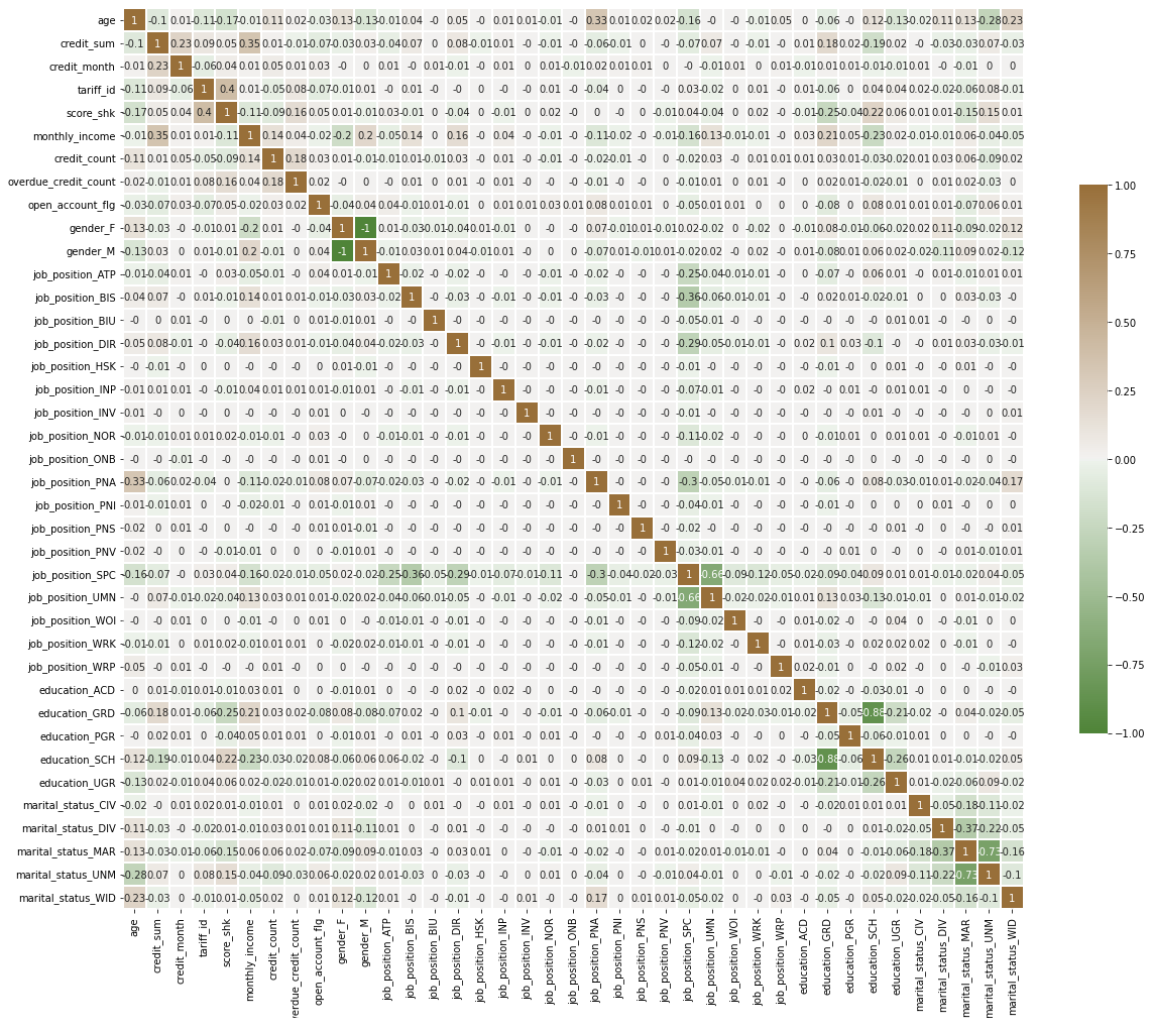
1) Корреляционный анализ

In [7]:

```

corr = data.corr().round(2)
f, ax = plt.subplots(figsize=(20, 20))
cmap = sns.diverging_palette(120, 50, as_cmap=True)
sns.heatmap(data=corr, cmap=cmap, annot=True, vmax=1.0, square=True, lin
plt.show()

```



In [8]:

```

print('Признаки, имеющие максимальную по модулю корреляцию с целевым при:
best_params = data.corr()['open_account_flg'].map(abs).sort_values(ascen

```

```
best_params = best_params[best_params.values > 0.02]
best_params
```

Out[8]:

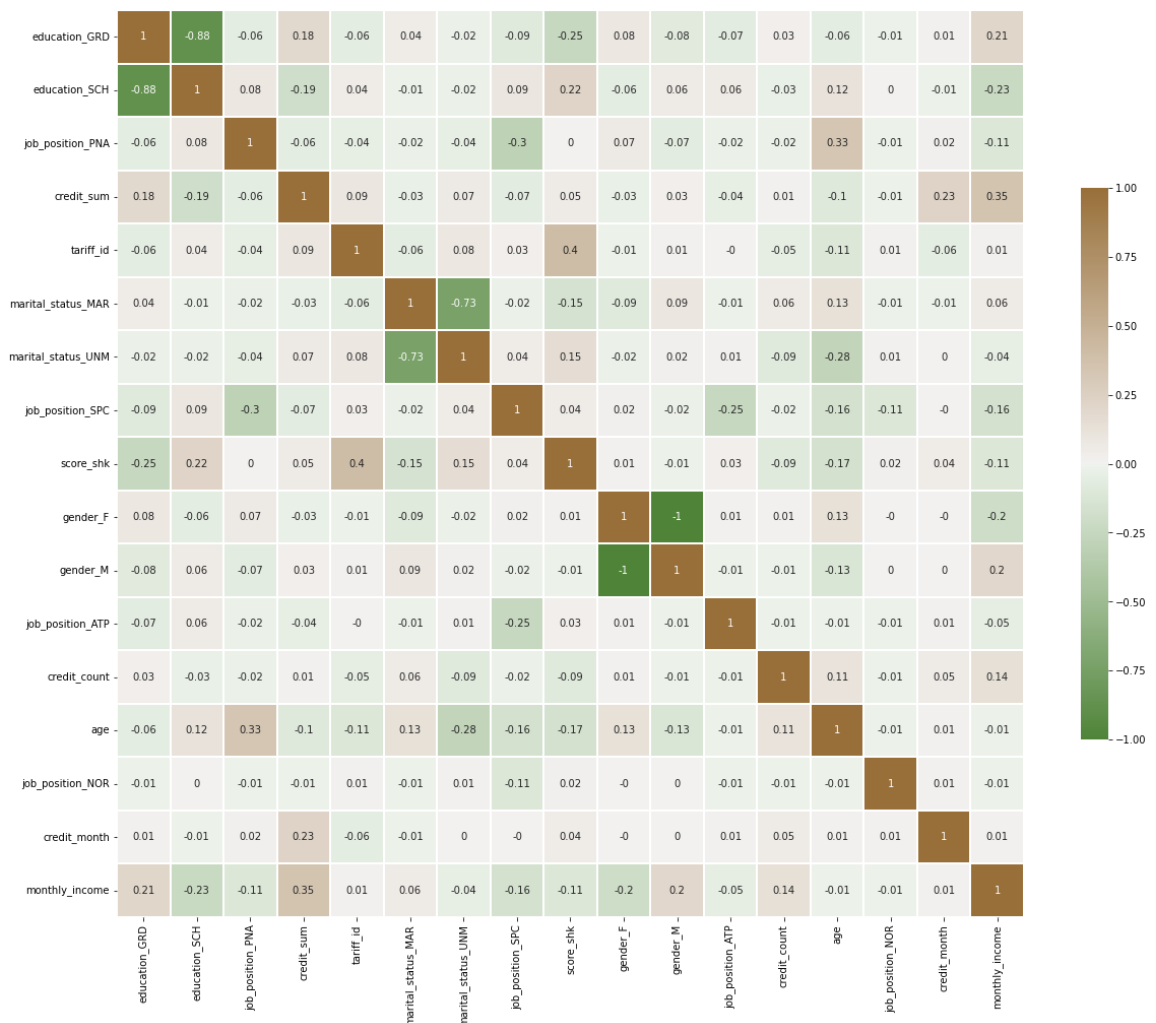
Признаки, имеющие максимальную по модулю корреляцию с целевым признаком

education_GRD	0.082371
education_SCH	0.078337
job_position_PNA	0.076889
credit_sum	0.072039
tariff_id	0.067346
marital_status_MAR	0.067112
marital_status_UNM	0.061312
job_position_SPC	0.049143
score_shk	0.048686
gender_F	0.044265
gender_M	0.044265
job_position_ATP	0.038288
credit_count	0.032374
age	0.031062
job_position_NOR	0.027320
credit_month	0.025809
monthly_income	0.023697


Name: open_account_flg, dtype: float64

In [9]:

```
corr = data[best_params.index].corr().round(2)
f, ax = plt.subplots(figsize=(20, 20))
cmap = sns.diverging_palette(120, 50, as_cmap=True)
sns.heatmap(data=corr, cmap=cmap, annot=True, vmax=1.0, square=True, lin
```



```
In [10]: plt.figure(figsize=(6, 8))
sns.heatmap(pd.DataFrame(data[np.append(best_params.index.values, 'open_
plt.show()
```



2) Разделение выборки на обучающую и тестовую

```
In [11]: data_best = data[best_params.index]
data_best.head()
```

```
Out[11]:
```

	education_GRD	education_SCH	job_position_PNA	credit_sum	tariff_id	marital_status_M
0	1	0	0	59998.00	1.6	
1	0	1	0	10889.00	1.1	
2	0	1	0	10728.00	1.1	
3	0	1	0	12009.09	1.1	
4	0	1	0	21229.00	1.1	

```
In [12]: y = data['open_account_flg']
#X = data.drop('open_account_flg', axis=1)
X = data_best
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.75)
x_train, x_test, y_train, y_test = train_test_split(x_train, y_train, te
```

3) Масштабирование данных

```
In [13]: scaler = MinMaxScaler().fit(x_train)
x_train = pd.DataFrame(scaler.transform(x_train), columns=x_train.columns)
x_test = pd.DataFrame(scaler.transform(x_test), columns=x_train.columns)
x_train.describe()
```

```
Out[13]:
```

	education_GRD	education_SCH	job_position_PNA	credit_sum	tariff_id	mari
count	29880.000000	29880.000000	29880.000000	29880.000000	29880.000000	
mean	0.425000	0.514759	0.023561	0.117340	0.345539	
std	0.494351	0.499790	0.151679	0.082275	0.252486	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.060249	0.106383	
50%	0.000000	1.000000	0.000000	0.092536	0.340426	
75%	1.000000	1.000000	0.000000	0.148270	0.638298	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

4) Модель №1: Случайный лес

```
In [82]: from sklearn.metrics import mean_absolute_error
from sklearn.metrics import median_absolute_error, r2_score

def print_metrics(y_test, y_pred):
    print(f"Precision: {precision_score(y_test, y_pred)}")
    print(f"F1-measure: {f1_score(y_test, y_pred)}")
    print('mean_absolute_error: {}'.format(round(mean_absolute_error(y_test, y_pred), 2)))
    print('median_absolute_error: {}'.format(round(median_absolute_error(y_test, y_pred), 2)))
    print('r2_score: {}'.format(round(r2_score(y_test, y_pred), 2)))
```

```
In [15]: print_metrics(y_test, RandomForestClassifier(random_state=17).fit(x_train, y_train))

Precision: 0.4737991266375546
F1-measure: 0.15617128463476068
```

Подбор гиперпараметров

```
In [90]: rf = RandomForestClassifier(random_state=17)
params = {'n_estimators': [100, 1000],
          'max_features': ['auto', 'sqrt'], 'min_samples_leaf': [1, 3, 5]}
grid_cv = GridSearchCV(estimator=rf, cv=5, param_grid=params, n_jobs=-1)
grid_cv.fit(x_train, y_train)
print(grid_cv.best_params_)

{'max_features': 'auto', 'min_samples_leaf': 5, 'n_estimators': 1000}
```

```
In [ ]: best_rf = grid_cv.best_estimator_
best_rf.fit(x_train, y_train)
y_pred_rf = best_rf.predict(x_test)
print_metrics(y_test, y_pred_rf)
```

5) Модель №2: Градиентный бустинг

```
In [16]: print_metrics(y_test, GradientBoostingClassifier(random_state=17).fit(x_train, y_train))
```

Accuracy: 0.8218803685772295
Precision: 0.5819672131147541
Recall: 0.06118052563550194
F1-measure: 0.11072124756335285

Подбор гиперпараметров

```
In [76]: gb = GradientBoostingClassifier(random_state=17)
params = {'n_estimators': [10, 50, 100, 200], 'min_samples_leaf': [1, 3, 5, 10]}
grid_cv = GridSearchCV(estimator=gb, cv=5, param_grid=params, n_jobs=-1)
grid_cv.fit(x_train, y_train)
print(grid_cv.best_params_)

{'min_samples_leaf': 5, 'n_estimators': 200}
```

```
In [77]: best_gb = grid_cv.best_estimator_
best_gb.fit(x_train, y_train)
y_pred_gb = best_gb.predict(x_test)
print_metrics(y_test, y_pred_gb)
```

Precision: 0.5709876543209876
F1-measure: 0.13988657844990549

6) Модель №3: Стекинг

```
In [24]: dataset = Dataset(x_train, y_train, x_test)
```

```
In [97]: from sklearn.ensemble import RandomForestClassifier, StackingClassifier,
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import SGDClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier

layer_one_estimators = [
    ('rf_1', RandomForestClassifier(n_estimators=10, random_state=17)),
    ('rf_3', GradientBoostingClassifier(n_estimators=100, random_state=17))
]

layer_two_estimators = [
    ('dt_2', DecisionTreeClassifier()),
    ('rf_2', RandomForestClassifier(n_estimators=10, random_state=17)),
    ('rf_4', GradientBoostingClassifier(n_estimators=100, random_state=17))
]

layer_two = StackingClassifier(estimators=layer_two_estimators, final_estimator=LogisticRegression())
# Create Final model by
clf = StackingClassifier(estimators=layer_one_estimators, final_estimator=layer_two)

#layer_2 = StackingClassifier(estimators=profiler_learners, final_estimator=LogisticRegression())
#layer_1 = StackingClassifier(estimators=base_learners, final_estimator=LogisticRegression())

clf.fit(x_train, y_train)
```

```
Out[97]: StackingClassifier(estimators=[('rf_1',
                                      RandomForestClassifier(n_estimators=10,
                                                              random_state=42)),
                                      ('rf_3',
                                      GradientBoostingClassifier(n_estimators=2
00))],
                                final_estimator=StackingClassifier(estimators=[('dt_
2',
                                                                              Decisi
onTreeClassifier()),
                                                                              ('rf_
2',
                                                                              Random
ForestClassifier(n_estimators=10,
random_state=42)),
                                                                              ('rf_
4',
                                                                              Gradie
ntBoostingClassifier(n_estimators=20))],
                                final_estimator=MLP
Classifier(random_state=1488)))
```

```
In [98]: print_metrics(y_test, clf.predict(x_test))
```

```
Precision: 0.6148148148148148
F1-measure: 0.12813585488228482
mean_absolute_error: 0.18
median_absolute_error: 0.0
r2_score: -0.19
```

Сравнение моделей

```
In [6]: print("Случайный лес")
print_metrics(y_test, y_pred_rf)

print("\nГрадиентный бустинг")
print_metrics(y_test, y_pred_gb)

print("\nСтекинг")
print_metrics(y_test, y_pred_stack)
```

Случайный лес

```
-----
--
NameError                                Traceback (most recent call las
t)
~\AppData\Local\Temp\ipykernel_12728\287930607.py in <module>
      1 print("Случайный лес")
----> 2 print_metrics(y_test, y_pred_rf)
      3
      4 print("\nГрадиентный бустинг")
      5 print_metrics(y_test, y_pred_gb)

NameError: name 'print_metrics' is not defined
```

```
In [ ]:
```